

Aiven Kafka

Quickstart with Python

by Yossi Drori | on 05 July 2022 | Technical How-To

In this blog post, I show you how to configure Kafka pipeline by using **Aiven** services and Python code

Inorder to check the pipeline, I will show how to build a client producer and to stream a stock market mockup data using Python

On the last step I will guide on how to Monitor the pipeline

Solution overview

The following architecture diagram-Figure 1-presents an overview of the solution.

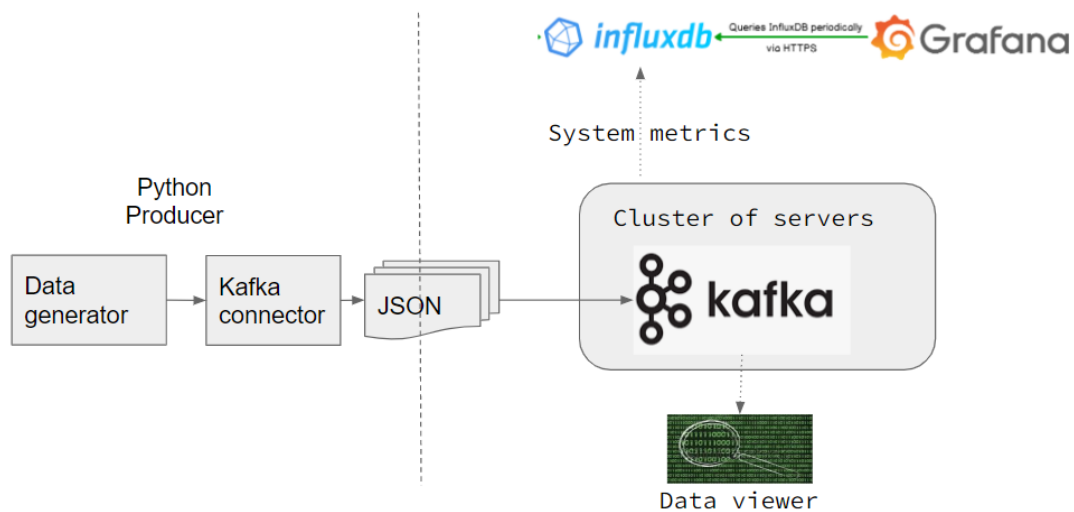


Figure 1: Solution architecture diagram

- In this example we can see 2 main components:
 - A. **The producer** - is responsible for generating data and sending it to the Kafka cluster
 - B. **The Kafka cluster** - is responsible for receiving the data

Prerequisites

- [Registering to Aiven](#)
- Python environment
 - [Installing faker](#)
 - [Installing kafka-python](#)
- Clone the github project: [Python Fake Data Producer for Apache Kafka](#)

Implement the solution

Setting up a Aiven Kafka service

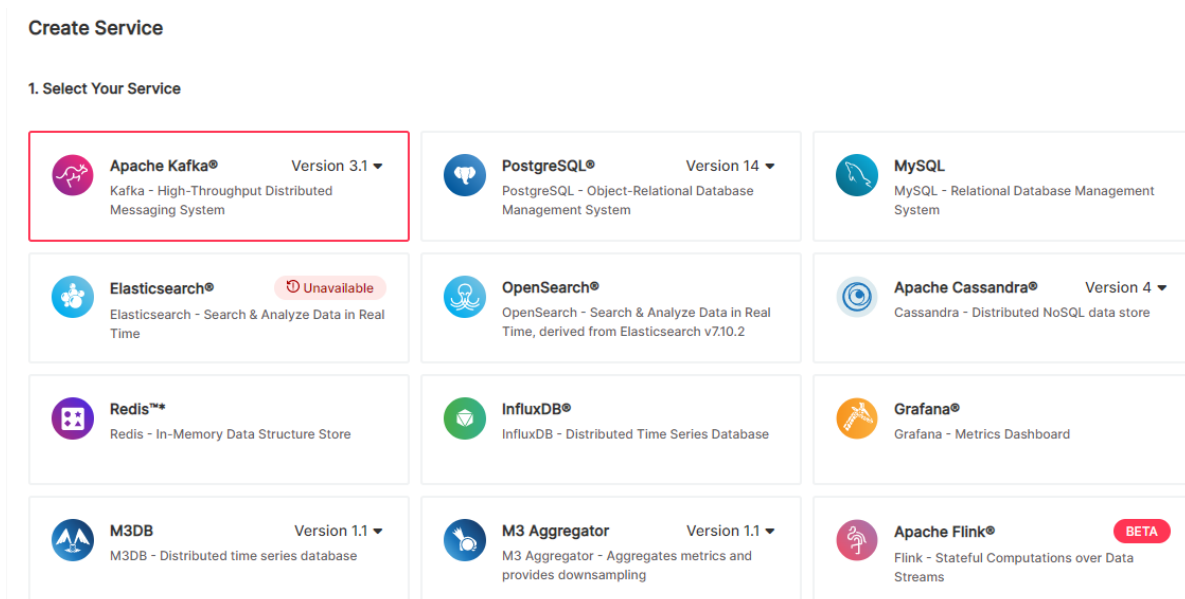


Figure 2: Choose the relevant service

From the list for all services we will first choose “Apache Kafka”

Kafka configuration

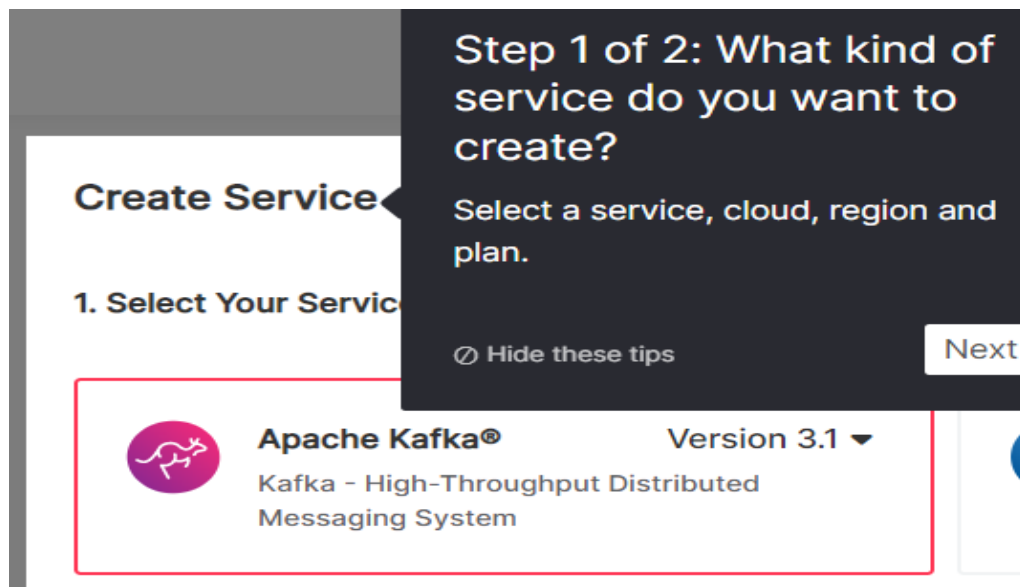



Figure 3: Kafka configuration tool tips

After selecting the desired service (Kafka in our case)
Most informative tool tips will guide you through on how to set it the best for your use case!

Selecting: Cloud service + Region

2. Select Service Cloud Provider



3. Select Service Cloud Region

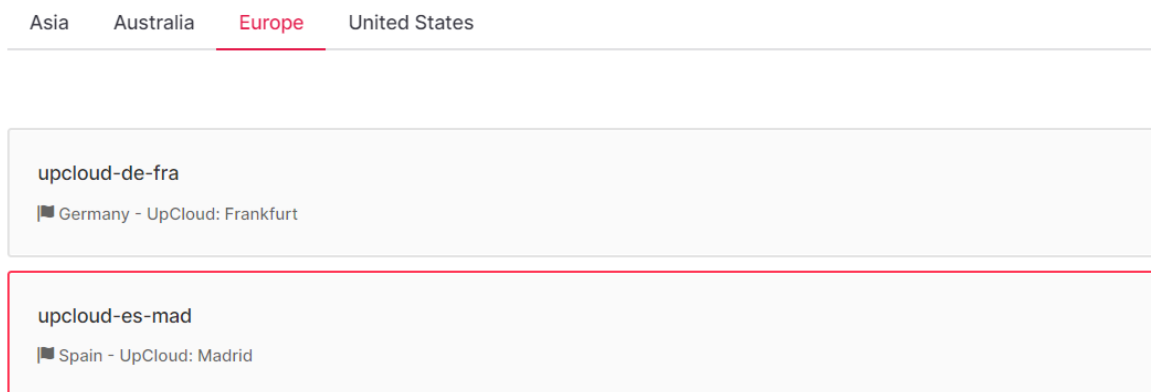


Figure 4: select Cloud service + Region

Flexible service plans

4. Select Service Plan

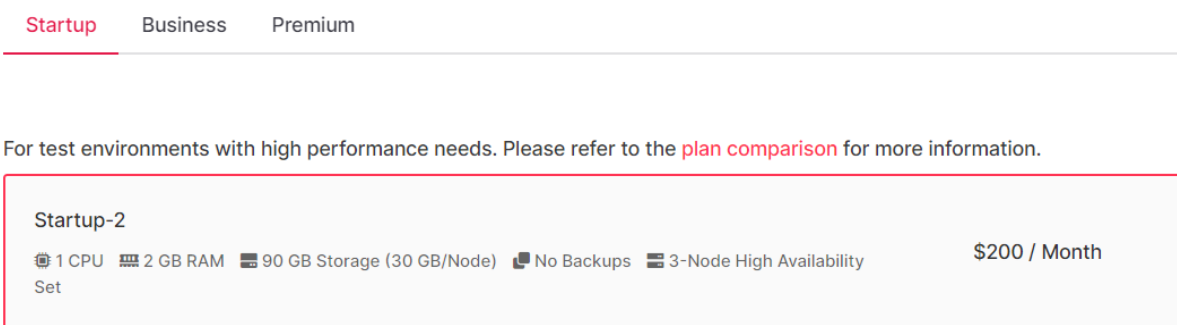



Figure 5: select a plan


In my case I'm using "Startup" since it's just a functionality test which requires light resources For "heavy lifting" you will consider "Business" / "Premium".

Advanced configurations

Edit advanced configuration ×


 INFO

Making advanced configuration changes may lead to your service restarting

kafka_authentication_methods.certificate* 


Synced

☒

kafka_authentication_methods.sasl* 

Synced


☐

kafka.auto_create_topics_enable* 

New

Not synced

☒



[+ Add configuration option](#)

Cancel

Save advanced configuration

Figure 6: Advanced configurations

Service setup summary

Service Summary

Name

kafka-1bd8f8a4

Service

 Apache Kafka 3.1

Cloud

 UpCloud

Region

Europe, Spain - UpCloud: Madrid

Plan

business-4

 2 CPU  4 GB RAM  600 GB Storage

 No Backups

 3-Node High Availability Set

Estimated Monthly Price*

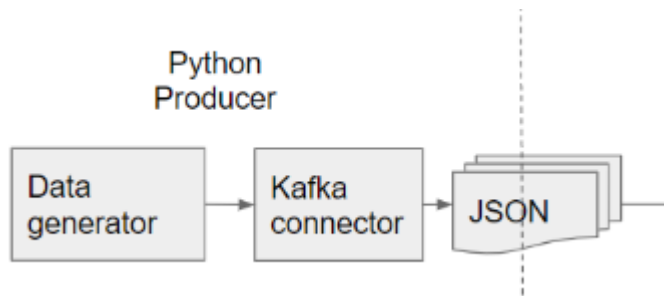
\$ 500 USD

*Estimated monthly price is based on 730 hours



Figure 6: Service setup summary

Data generator and producer



See - Figure 1

Data generator

In order to generate the mockup stock tickets data I have used [Faker Python lib](#)

In my example I'm generating stock tickets mockup data

The file [stockproducer.py](#) is implementing the code for the mockup data:

- **Key**

```
key = str(uuid.uuid4())
```

- **stock_name**

```
def stock_name(self):  
    return random.choice(StockNames)
```

- **stock_value**

```
def stock_value(self, stockname):  
    indexStock = StockNames.index(stockname)  
    currentval = StockCurrentValues[indexStock]  
    goesup = 1  
    if random.random() > StockUpProb[indexStock]:  
        goesup = -1  
    nextval = currentval + random.random() * ChangeAmount * goesup  
    StockCurrentValues[indexStock] = nextval
```

- **timestamp** - Is actually the current time in ISO 8601 format

```
'timestamp': datetime.now().isoformat()
```

An example for generated event:

key: cd4fab54-6f6e-4f30-ac93-7e1941079357

message:

```
{  
    'stock_name': 'Indiana Jeans'  
    , 'stock_value': 26.568693883965526  
    , 'timestamp': '2022-07-05T23:23:58.236269'  
}
```

Sending the data into Kafka cluster

On the [main.py](#) I've implemented the interaction with the Kafka cluster by using the "KafkaProducer"

```
from kafka import KafkaProducer
```

To run the Python producer

```
python main.py
```

--cert-folder *Stock*

The folder with the credential files

--host *kafka-stock-tickets-drori-a69e.aivencloud.com*

The Kafka service URL

--port *28003*

--topic-name *stock-tickets*

The Kafka topic name (it will be created dynamically on the Kafka side)

--nr-messages *1000*

Number of messages to generate

--max-waiting-time *1*

Maximum time between messages

```
Sleeping for...0.095s
Sending...key:bfbf5387-4c6e-4cb3-ab96-ebfc41d7f9bf message: {'stock_name': 'Indiana Jeans', 'stock_value': 27.043917298293493, 'timestamp': '2022-07-06T07:23:57.689540'}
Sleeping for...0.1654s
Sending...key:3d19855c-3890-4c9b-bf76-fdef234e2085 message: {'stock_name': 'Jurassic Pork', 'stock_value': 19.41300276008913, 'timestamp': '2022-07-06T07:23:57.861380'}
Sleeping for...0.067s
Sending...key:f99a191d-244e-426c-867e-f15922c09022 message: {'stock_name': 'Jurassic Pork', 'stock_value': 19.19835195223107, 'timestamp': '2022-07-06T07:23:57.939483'}
Sleeping for...0.289s
Sending...key:cd4fab54-6f6e-4f30-ac93-7e1941079357 message: {'stock_name': 'Indiana Jeans', 'stock_value': 26.568693883965526, 'timestamp': '2022-07-06T07:23:58.236269'}
Sleeping for...0.0951s
```

Observability and Monitoring

Browsing to the kafka topic.
The generated data includes stock tickets values

When browsing to the topic we can see the JSON steamed data

We can see 2 sections: “key” and “value”

Key: UUID

VALUE: {“stock_name”, “stock_value”, “timestamp”}

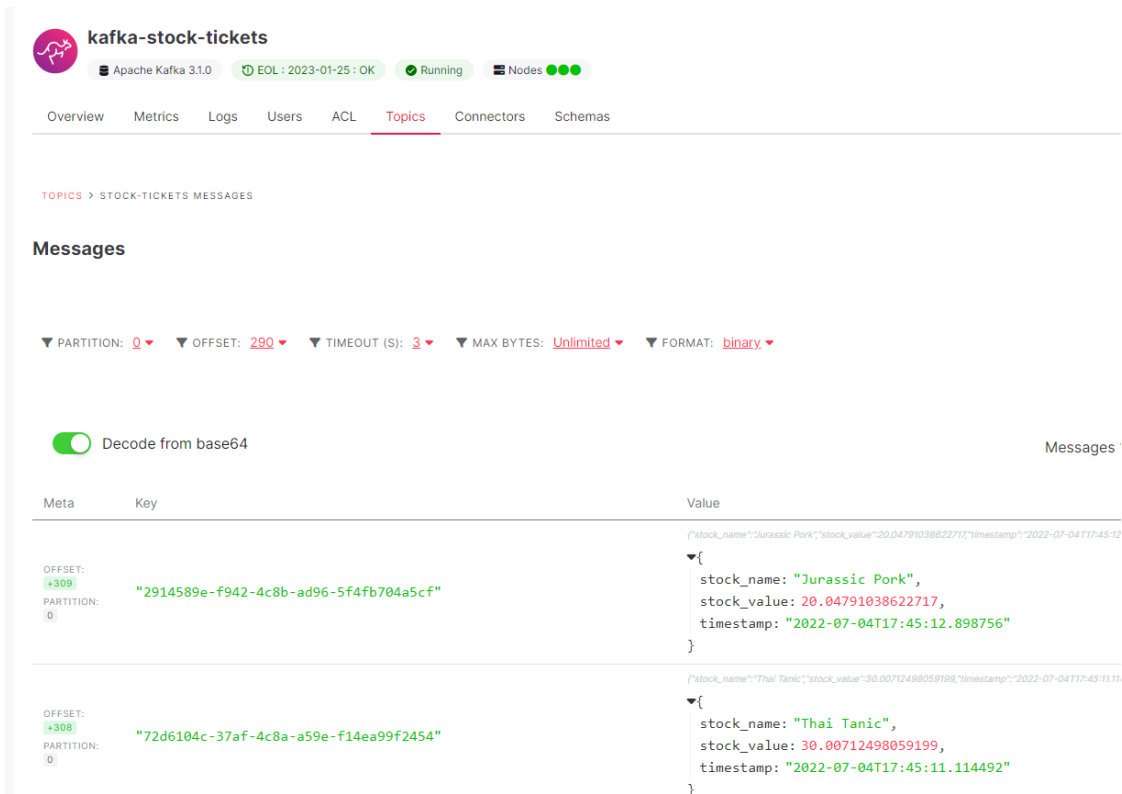


Figure 7: Browsing to the data in the pipeline

This view is available from: Kafka service view => Topics => Topic => Fetch Messages

Monitoring system matrix



Service	
	influx-e4e010d InfluxDB • Powered off
	grafana-stock-tickets-pipeline Grafana • Powered off

Figure 8: Adding services to enable system monitoring

Inorder to have a system monitoring I've created 2 additional services:

- 1. **Influx-db** - A database that connects to the Kafka cluster and collects the system matrices
- 2. **Grafana** - A visualizer that is connected to the **Influx-db** and can present all the system matrices in a dashboard



See - Figure 1

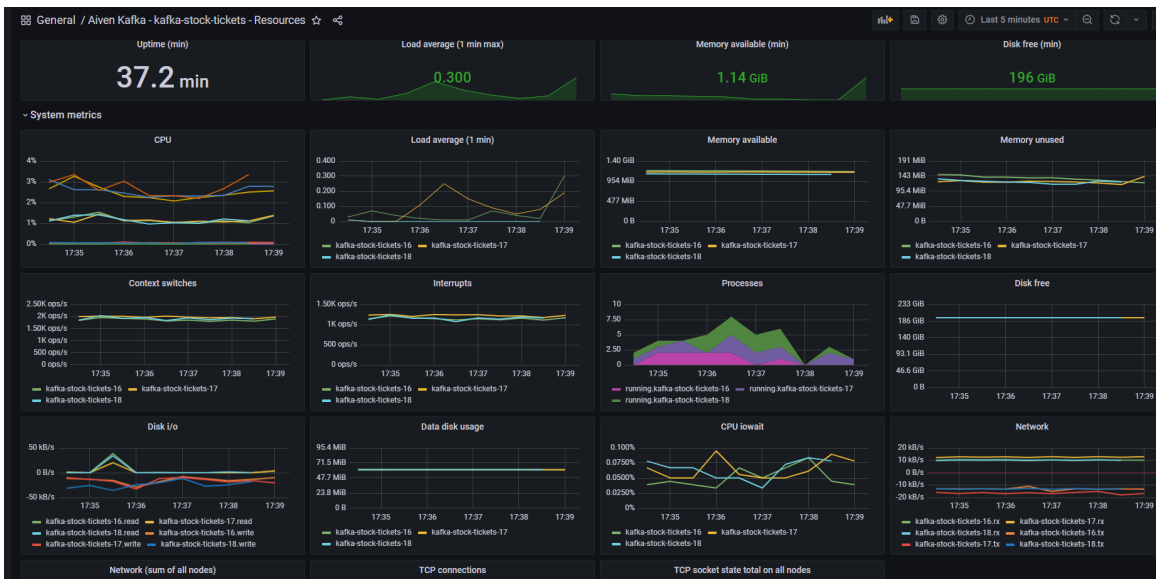


Figure 8: Using Grafana to monitor the system matrices

Wrapping up

As I have shown in this document, by following a simple set of instructions you can easily set up an event driven pipeline.

Working with Aiven console is very simple and have a lot of tool tips to guide you through

Since the service is a “Pay as you go” don’t forget that you can stop the services when you don’t need it up and running.

For any question or feedback please contact me or [our support team](#)

Email: drori.yossi@gmail.com